

Localization for Autonomous Driving

Brief tutorial to have fun with localization

Prepared by Oussama Code by Claude.ai

January 2025

I, Oussama EL HAMZAOUI confirm that the work presented in this report is my own (with the help of Claude.ai). Where information has been derived from other sources, I confirm that this has been indicated in the report.

Abstract

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Nam et turpis gravida, lacinia ante sit amet, sollicitudin erat. Aliquam efficitur vehicula leo sed condimentum. Phasellus lobortis eros vitae rutrum egestas. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia Curae; Donec at urna imperdiet, vulputate orci eu, sollicitudin leo. Donec nec dui sagittis, malesuada erat eget, vulputate tellus. Nam ullamcorper efficitur iaculis. Mauris eu vehicula nibh. In lectus turpis, tempor at felis a, egestas fermentum massa.

Acknowledgements

Interdum et malesuada fames ac ante ipsum primis in faucibus. Aliquam congue fermentum ante, semper porta nisl consectetur ut. Duis ornare sit amet dui ac faucibus. Phasellus ullamcorper leo vitae arcu ultricies cursus. Duis tristique lacus eget metus bibendum, at dapibus ante malesuada. In dictum nulla nec porta varius. Fusce et elit eget sapien fringilla maximus in sit amet dui.

Mauris eget blandit nisi, faucibus imperdiet odio. Suspendisse blandit dolor sed tellus venenatis, venenatis fringilla turpis pretium. Donec pharetra arcu vitae euismod tincidunt. Morbi ut turpis volutpat, ultrices felis non, finibus justo. Proin convallis accumsan sem ac vulputate. Sed rhoncus ipsum eu urna placerat, sed rhoncus erat facilisis. Praesent vitae vestibulum dui. Proin interdum tellus ac velit varius, sed finibus turpis placerat.

Contents

Abstract	i
Acknowledgements	ii
Abbreviations	
Overview	
Course chapters	
Objectives	1
1 Introduction to Localization	2
1.1 Problem Statement and Motivation	2
1.2 Types of Localization Problems	2
1.3 Sensor Types and Characteristics	3
1.4 Sources of Uncertainty	3
1.5 The Role of Probability Theory	3
2 Probability Theory Foundations	5
2.1 Random variables and probability distributions	5
2.1.1 Random Variables	5
2.1.2 Probability Distributions	5
2.2 Bayes theorem	6
2.3 Conditional probability	7
2.4 Markov assumption	8
2.5 Joint and marginal probabilities	9
2.6 Gaussian distributions	11
3 Bayesian Filtering Framework	13
3.1 Recursive state estimation	13
3.2 Prediction step (motion model)	14
3.3 Update step (measurement model)	15
3.4 Chapman-Kolmogorov equation	17
3.5 Bayes filter algorithm	18
3.6 Linear vs nonlinear systems	20
3.6.1 Linear Systems	20
3.6.2 Nonlinear Systems	21

4	Kalman Filtering	22
4.1	Linear Kalman Filter	22
4.1.1	System model and assumptions	22
4.1.2	Prediction equations	23
4.1.3	Update equations	24
4.1.4	Uncertainty propagation	26
4.2	Extended Kalman Filter (EKF)	27
4.2.1	Linearization process	27
4.2.2	Jacobian matrices	30
4.2.3	Algorithm implementation	31
4.3	Unscented Kalman Filter (UKF)	39
4.3.1	Sigma points	39
4.3.2	Unscented transform	40
4.3.3	Algorithm implementation	42
5	Particle Filtering	49
6	Advanced Topics	50
7	MATLAB Implementation: Kalman Filter	51
8	MATLAB Implementation: Particle Filter	52
9	Practical Applications	53
10	Project Work	54

List of Figures

List of Tables

2.1	Joint and Marginal Probabilities of Cold and Rainy Days	10
-----	---	----

Abbreviations

ADAS	Advanced D river A ssistance S ystems
CKF	Cubature K alman F ilter
EKF	Extended K alman F ilter
GNSS	Global Navigation Satellite S ystem
HD	H igh D efinition (as in HD maps)
IMU	Inertial Measurement U nit
KF	K alman F ilter
LIDAR	Light D etection and R anging
MCL	Monte C arlo L ocalization
PF	P article F ilter
RADAR	Radio D etection and R anging
RBPF	Rao-Blackwellized P article F ilter
SLAM	Simultaneous L ocalization and M apping
UKF	Unscented K alman F ilter
V2X	Vehicle-to-Everything Communication

Overview

Course chapters

Section to be deleted after completion of the course

- ☒ Introduction to Localization
 - ☒ Problem statement and motivation
 - ☒ Types of localization problems
 - ☒ State estimation challenges
 - ☒ Sensor types and characteristics
 - ☒ Sources of uncertainty in robotics
- ☒ Probability Theory Foundations
 - ☒ Random variables and probability distributions
 - ☒ Bayes' theorem
 - ☒ Conditional probability
 - ☒ Markov assumption
 - ☒ Joint and marginal probabilities
 - ☒ Gaussian distributions
- ☒ Bayesian Filtering Framework
 - ☒ Recursive state estimation
 - ☒ Prediction step (motion model)
 - ☒ Update step (measurement model)
 - ☒ Chapman-Kolmogorov equation
 - ☒ Bayes filter algorithm
 - ☒ Linear vs nonlinear systems
- [-] Kalman Filtering
 - ☒ Linear Kalman Filter
 - ☒ System model and assumptions
 - ☒ Prediction equations
 - ☒ Update equations
 - ☒ Uncertainty propagation
 - ☒ Extended Kalman Filter (EKF)
 - ☒ Linearization process
 - ☒ Jacobian matrices
 - ☒ Algorithm implementation
 - [-] Unscented Kalman Filter (UKF)
 - ☒ Sigma points
 - ☒ Unscented transform
 - ☒ Algorithm implementation
- ☐ Particle Filtering
 - ☐ Monte Carlo methods

- ☐ Importance sampling
- ☐ Particle representation
- ☐ Sequential Importance Sampling (SIS)
- ☐ Resampling techniques
- ☐ Sample degeneracy and impoverishment
- ☐ Adaptive particle filtering
- ☐ Advanced Topics
 - ☐ Multi-hypothesis tracking
 - ☐ SLAM basics
 - ☐ Sensor fusion techniques
 - ☐ Loop closure
 - ☐ Global vs local localization
- ☐ MATLAB Implementation: Kalman Filter
 - ☐ Linear KF implementation
 - ☐ State prediction
 - ☐ Measurement update
 - ☐ Covariance propagation
 - ☐ EKF implementation
 - ☐ System modeling
 - ☐ Jacobian computation
 - ☐ Filter implementation
 - ☐ Visualization and analysis
 - ☐ Performance evaluation
- ☐ MATLAB Implementation: Particle Filter
 - ☐ Basic PF framework
 - ☐ Particle initialization
 - ☐ Motion model implementation
 - ☐ Measurement model
 - ☐ Weight computation
 - ☐ Resampling implementation
 - ☐ Visualization tools
 - ☐ Performance metrics
- ☐ Practical Applications
 - ☐ Vehicle localization case studies
 - ☐ Robot navigation examples
 - ☐ Integration with mapping
 - ☐ Real-world challenges
 - ☐ Best practices and optimization
- ☐ Project Work
 - ☐ Implementation exercises
 - ☐ Real dataset analysis
 - ☐ Performance comparison of different filters
 - ☐ Parameter tuning
 - ☐ Documentation and presentation

Objectives

Chapter 1

Introduction to Localization

- ☒ Introduction to Localization
 - ☒ Problem statement and motivation
 - ☒ Types of localization problems
 - ☒ State estimation challenges
 - ☒ Sensor types and characteristics
 - ☒ Sources of uncertainty in robotics

1.1 Problem Statement and Motivation

In robotics and autonomous systems, localization addresses a fundamental question: “Where am I?” This seemingly simple question underlies many complex challenges in autonomous navigation. Imagine waking up in an unfamiliar room – you would use visual cues, memory, and perhaps a map to determine your location. Robots face a similar challenge, but must solve it using sensors and algorithms rather than human intuition.

Localization serves as the cornerstone of autonomous navigation. Without accurate knowledge of its position, a robot cannot effectively plan paths, avoid obstacles, or complete assigned tasks. This becomes particularly critical in applications like autonomous vehicles, where position errors of even a few centimeters can have serious consequences.

1.2 Types of Localization Problems

We can categorize localization problems based on their initial conditions and objectives:

Position Tracking represents the simplest case, where we know the initial position and need to maintain an accurate estimate as the robot moves. Think of using GPS in your car – you start from a known location and track your movement.

Global Localization presents a more challenging scenario where the initial position is unknown. The robot must determine its position from scratch using available sensor information and a map. This is analogous to opening a ride-sharing app in an unfamiliar city and waiting for it to locate you.

The Kidnapped Robot Problem is the most challenging variant, where a well-localized robot is suddenly transported to an unknown location. While this may seem artificial, it tests a system's ability to recover from catastrophic failures or sensor malfunctions.

1.3 Sensor Types and Characteristics

Localization systems typically rely on multiple sensor types, each with distinct advantages and limitations:

- Proprioceptive Sensors measure internal state changes, such as wheel encoders that track rotation or inertial measurement units (IMUs) that detect acceleration and angular velocity. While these sensors provide high-frequency updates, they suffer from cumulative errors through a process called dead reckoning.
- Exteroceptive Sensors observe the external environment. These include:
 - LIDAR (Light Detection and Ranging) which creates detailed 3D scans of surroundings
 - Cameras that provide rich visual information but require sophisticated processing
 - RADAR which offers reliable distance measurements even in adverse weather
 - GNSS (Global Navigation Satellite System) which provides absolute position but may suffer from urban canyon effects and multipath errors

1.4 Sources of Uncertainty

Understanding uncertainty is crucial for robust localization. Several factors contribute to localization uncertainty:

Motion Uncertainty arises from imperfect robot control and environmental interactions. When a robot moves, wheel slippage, uneven terrain, and mechanical play all introduce errors between commanded and actual motion.

Measurement Uncertainty stems from sensor limitations and noise. For example, LIDAR measurements might be affected by reflective surfaces, while camera images can be distorted by varying lighting conditions.

Environmental Uncertainty relates to the dynamic nature of the real world. Moving objects, changing weather conditions, and modifications to the environment can all affect localization accuracy.

Model Uncertainty comes from our simplified representations of complex physical systems. Our mathematical models of robot motion and sensor behavior are approximations that introduce additional uncertainty.

1.5 The Role of Probability Theory

Given these uncertainties, deterministic approaches to localization often fail in real-world conditions. This necessitates a probabilistic framework that can: - Represent and propa-

gate uncertainty through mathematical models - Fuse information from multiple, imperfect sensors - Handle conflicting measurements and outliers - Provide confidence estimates along with position estimates

This probabilistic approach leads us naturally to the Bayesian filtering framework, which we'll explore in subsequent chapters. The framework provides a mathematical foundation for combining prior knowledge, motion predictions, and sensor measurements to maintain an estimate of the robot's position over time.

Understanding these foundational concepts is crucial as we progress to more advanced topics in localization. The challenges and considerations introduced here will inform our discussion of specific algorithms and implementations throughout the course.

Would you like me to elaborate on any of these sections or move on to the probability theory foundations?

Chapter 2

Probability Theory Foundations

2.1 Random variables and probability distributions

2.1.1 Random Variables

A random variable is a mathematical way to describe outcomes of a random process. Think of it as a function that assigns a numerical value to each possible outcome of an experiment or observation.

Let's consider a practical example from robotics: imagine a robot's sensor measuring the distance to a wall. Even when the robot and wall are stationary, repeated measurements might give slightly different values due to sensor noise. Each measurement is a realization of a random variable that we could call "measured distance."

Random variables come in two main types:

Discrete random variables can only take specific, countable values. For instance, if we count the number of landmarks a robot sees in its field of view, this would be a discrete random variable - we can only see 0, 1, 2, or some whole number of landmarks.

Continuous random variables can take any value within a continuous range. Most sensor measurements in robotics are continuous random variables. Our distance sensor example could theoretically return any real number within its measurement range.

2.1.2 Probability Distributions

A probability distribution describes how likely each possible value of a random variable is to occur. It tells us the complete story of the random variable's behavior.

For discrete random variables, we use a Probability Mass Function (PMF). The PMF gives the probability of each possible value directly. For example, if we're counting landmarks:

- $P(X = 0) = 0.1$ (10% chance of seeing no landmarks)
- $P(X = 1) = 0.3$ (30% chance of seeing exactly one landmark)
- $P(X = 2) = 0.4$ (40% chance of seeing exactly two landmarks)

And so on...

For continuous random variables, we use a Probability Density Function (PDF). The PDF works differently because with continuous variables, the probability of getting any exact value is actually zero! Instead, the PDF gives us the relative likelihood of values occurring, and we integrate it over ranges to get probabilities.

The most important continuous probability distribution in robotics is the Gaussian (or Normal) distribution. It's defined by two parameters:

- (mu): the mean, representing the central value
- (sigma): the standard deviation, representing the spread

The Gaussian distribution appears naturally in many robotics scenarios because of the Central Limit Theorem. When many small random effects add up - like multiple sources of sensor noise - their combined effect tends to follow a Gaussian distribution.

In the context of localization, probability distributions help us represent:

1. The robot's belief about its position (often as a Gaussian in simple cases)
2. Uncertainty in sensor measurements
3. Noise in motion commands and their execution
4. The likelihood of different measurements given a particular position

Understanding these distributions is crucial because localization algorithms like Kalman filters and particle filters essentially manipulate these probability distributions to maintain and update the robot's position estimate over time.

2.2 Bayes theorem

Let me explain Bayes' theorem in a way that will make intuitive sense, starting with a simple example and then building up to its use in robotics.

Imagine you're a robot in a room, and you have a simple distance sensor. Sometimes your sensor shows a reading of 2 meters, but you're not sure if you're actually 2 meters from a wall or if your sensor is giving you a wrong reading.

To understand Bayes' theorem, let's break this situation down into pieces:

First, let's define what we know: - You might be 2 meters from a wall (we'll call this your "position") - Your sensor gives you a measurement of 2 meters (we'll call this your "measurement")

Now, what Bayes' theorem helps us figure out is: Given that your sensor reads 2 meters, what's the probability that you're actually 2 meters from the wall?

Here's the magic formula (don't worry, we'll break it down):

$$P(\text{position} \mid \text{measurement}) = P(\text{measurement} \mid \text{position}) \times P(\text{position}) / P(\text{measurement})$$

Let's understand each piece:

1. $P(\text{position} \mid \text{measurement})$ is what we want to know: the probability of being at a position, given our sensor measurement. This is called the "posterior probability."

2. $P(\text{measurement} \mid \text{position})$ is how likely we are to get this measurement if we really are at that position. We know our sensor isn't perfect - maybe it's 90% accurate when we're actually at 2 meters. This is called the "likelihood."
3. $P(\text{position})$ is what we believed about our position before taking the measurement. Maybe based on our last estimate, we thought there was a 70% chance we were at 2 meters. This is called the "prior probability."
4. $P(\text{measurement})$ is how likely we are to get this measurement in general. Think of it as a normalizing factor that makes all our probabilities add up to 100%.

Let's put some numbers in: - If our sensor is 90% accurate: $P(\text{measurement} \mid \text{position}) = 0.9$ - If we thought we were probably at 2m: $P(\text{position}) = 0.7$ - Let's say $P(\text{measurement}) = 0.8$ (this is calculated considering all possibilities)

Then:

$$P(\text{position} \mid \text{measurement}) = 0.9 \times 0.7 / 0.8 = 0.79$$

This tells us that after getting the measurement, we're 79% confident about our position - more confident than our prior belief of 70%!

The beautiful thing about Bayes' theorem is that it gives us a formal way to: 1. Start with what we believe (prior) 2. Consider new evidence (likelihood) 3. Update our belief (posterior)

In robotics, we use this process continuously. Every time we: - Move (this changes our prior belief) - Take a measurement (this gives us new evidence) - We use Bayes' theorem to update our belief about where we are

This is the foundation of probabilistic robotics and the basis for algorithms like Kalman filters and particle filters, which we'll explore later.

2.3 Conditional probability

Think of probability as measuring how likely something is to happen. Now, conditional probability takes this a step further by asking: "How likely is this event to happen, given that we already know something else has happened?"

The formal notation for conditional probability is $P(A|B)$, which reads as "the probability of A given B." Mathematically, it's expressed as:

$$P(A|B) = P(A \cap B) / P(B)$$

Let's break this down with a real-world example. Imagine we have a deck of 52 playing cards, and we want to know the probability of drawing a king, given that we've already drawn a red card.

To solve this:

1. First, we identify what we know: we've drawn a red card (this is our condition B)
2. We want to find the probability of having a king among these red cards (this is our event A)

3. $P(B)$ = probability of drawing a red card = $26/52 = 1/2$
4. $P(A \cap B)$ = probability of drawing a red king = $2/52 = 1/26$
5. Therefore, $P(A|B) = (2/52)/(26/52) = 2/26 = 1/13$

This shows us something interesting: while the probability of drawing a king from the full deck is $4/52$ (about 0.077), the probability of drawing a king given that we know the card is red is $1/13$ (about 0.077). In this case, knowing the card is red didn't change the probability of it being a king, because kings are evenly distributed between red and black cards.

This leads us to an important concept: independence. If knowing one event doesn't affect the probability of another event, we say these events are independent. In such cases, $P(A|B) = P(A)$. However, in many real-world scenarios, events are dependent, and conditional probability helps us account for this dependency.

Consider a medical example: the probability of having a certain disease might be 1% in the general population, but if we know a person has a specific symptom, the conditional probability of having the disease given this symptom might be much higher, say 30%. This is why doctors use symptoms to update their diagnostic probabilities.

2.4 Markov assumption

The Markov assumption, also known as the Markov property, is a fundamental concept in probability theory that helps us model complex sequences of events in a manageable way. Let me break this down step by step.

The core idea of the Markov assumption is that the future state of a system depends only on its present state, not on its past states. In probability terms, this means that if we want to predict what happens next, we only need to know what's happening right now, not the entire history of what happened before.

To understand this more concretely, imagine you're watching the weather. A pure Markov process would say that tomorrow's weather only depends on today's weather, not on what the weather was like last week or last month. While this might seem like an oversimplification (and in reality, weather patterns are more complex), this assumption often proves surprisingly useful in many real-world applications.

Let's express this mathematically. For a sequence of events $X_1, X_2, X_3, \dots, X_n$ the Markov property states that:

$$P(X_{n+1}|X_n, X_{n-1}, \dots, X_1) = P(X_{n+1}|X_n)$$

This equation tells us that the probability of the next state (X_{n+1}) given all previous states is equal to the probability of the next state given just the current state (X_n). This dramatically simplifies our calculations while still capturing many important patterns in real-world processes.

Think of it like playing a game of chess. While each position arose from a long sequence of moves, a player really only needs to look at the current board position to

decide their next move. The specific sequence of moves that led to this position, while interesting historically, isn't directly relevant to choosing the next best move.

The Markov assumption is particularly powerful because it allows us to build practical models of complex systems. It's used in:

1. Natural Language Processing - In simple language models, the probability of the next word might depend only on the current word (or last few words), not the entire sentence history.
2. Financial Markets - Some basic models assume that tomorrow's stock price depends only on today's price, not the entire price history.
3. Biology - Gene sequences can be modeled using Markov chains, where each base pair depends only on the previous few pairs.
4. Machine Learning - Hidden Markov Models use this property to model sequential data in a computationally efficient way.

It's important to note that the Markov assumption comes in different "orders." What I've described is a first-order Markov process, where we only look at the immediate previous state. In a second-order Markov process, we look at the last two states, and so on. Higher-order Markov processes can capture more complex dependencies but require more computational resources.

This assumption, while powerful, isn't always perfectly accurate in real-world situations. Many processes have longer-term dependencies that a simple Markov model might miss. However, the simplification it provides often outweighs these limitations, making it an invaluable tool in probability theory and its applications.

2.5 Joint and marginal probabilities

Let me explain joint and marginal probabilities through an intuitive progression, starting with the fundamentals and building up to how they work together.

Joint probability represents the likelihood of two (or more) events occurring together. We write this as $P(A, B)$ or $P(A \cap B)$, which reads as "the probability of A and B happening." Think of it as the overlap in a Venn diagram - the space where both events occur simultaneously.

Let's make this concrete with an example. Imagine we're looking at weather data for a year. Let's consider two events:

- Event A : It's a cold day (temperature below $50^{\circ}F$)
- Event B : It's a rainy day

The joint probability $P(A, B)$ would tell us the probability that a randomly chosen day is both cold AND rainy. If 73 days out of 365 were both cold and rainy, the joint probability would be $73/365 = 0.2$, or 20%.

Now, marginal probability is what we get when we're interested in the probability of just one event, regardless of what happens with other events. It's called "marginal"

because historically, these probabilities were written in the margins of probability tables. If we look at our weather example:

- $P(A)$: The probability of a cold day, regardless of rain
- $P(B)$: The probability of a rainy day, regardless of temperature

Here's where these concepts connect: marginal probabilities can be calculated by summing up joint probabilities. Mathematically:

$$P(A) = P(A, B) + P(A, \text{not} B)$$

In our weather example, if:

- 73 days were cold and rainy: $P(A, B) = 0.2$
- 109 days were cold and not rainy: $P(A, \text{not} B) = 0.3$

Then the marginal probability of a cold day $P(A) = 0.2 + 0.3 = 0.5$, or 50% of days.

We can visualize this with a probability table:

Table 2.1: Joint and Marginal Probabilities of Cold and Rainy Days. The sum of all joint probabilities equals 1, representing all possible weather combinations.

	Rainy (B)	Not Rainy	Marginal
Cold (A)	0.2	0.3	0.5
Not Cold	0.1	0.4	0.5
Marginal	0.3	0.7	1.0

Each cell shows a joint probability, and the margins show marginal probabilities. Notice how the marginals sum up the joint probabilities in their respective rows or columns.

This relationship between joint and marginal probabilities becomes especially important when working with conditional probabilities. Remember our earlier discussion about conditional probability? We can express it using joint and marginal probabilities:

$$P(A|B) = P(A, B)/P(B)$$

This shows how these concepts are deeply interconnected: joint probabilities help us calculate marginal probabilities, which in turn help us work with conditional probabilities.

Understanding these relationships is crucial in many real-world applications, from medical diagnosis (where we might look at joint probabilities of symptoms and diseases) to market analysis (examining the relationship between customer demographics and purchasing behaviors).

2.6 Gaussian distributions

We will dive into Gaussian distributions, also known as normal distributions, by building up from their fundamental characteristics to their broader significance in probability theory.

The Gaussian distribution is perhaps nature's most remarkable probability distribution. Picture a perfectly symmetrical bell-shaped curve - this is the visual signature of a Gaussian distribution. But why is this shape so special and ubiquitous?

At its core, a Gaussian distribution is defined by two key parameters:

1. The mean (μ) - the center point of the distribution, representing the average value
2. The standard deviation (σ) - which determines how spread out the values are from the mean

The mathematical formula for a Gaussian distribution is:

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$$

While this formula might look intimidating, we can understand its behavior through some key properties:

First, consider symmetry. The distribution is perfectly symmetrical around its mean, meaning values equally far above and below the mean are equally likely. This reflects many natural phenomena - for instance, human height variations around the average height.

Second, the “68 – 95 – 99.7 rule” tells us something remarkable about how probability is distributed:

- About 68% of values fall within one standard deviation of the mean
- About 95% fall within two standard deviations
- About 99.7% fall within three standard deviations

This leads us to an important insight: extreme values become exponentially less likely as we move away from the mean. Think about human height - while it's common to meet someone 2 inches taller than average, it's extremely rare to meet someone 12 inches taller.

The Gaussian distribution emerges naturally in many situations due to the Central Limit Theorem, which tells us that when we add up many independent random variables, their sum tends to follow a Gaussian distribution, regardless of the underlying distributions of the individual variables. This explains why we see Gaussian distributions so often in nature - many natural phenomena are the result of multiple small, independent effects adding together.

Let's consider some practical examples:

- Measurement Error: When scientists make repeated measurements of the same quantity, the errors typically follow a Gaussian distribution
- Financial Returns: Daily stock market returns often approximately follow a Gaussian distribution

- Biological Variations: Things like height, weight, and blood pressure in a population often follow roughly Gaussian distributions

Understanding Gaussian distributions is crucial for:

1. Statistical Testing: Many statistical tests assume underlying Gaussian distributions
2. Quality Control: Manufacturing processes often use Gaussian assumptions to set acceptable tolerance limits
3. Risk Analysis: Financial models frequently use Gaussian distributions to model market behavior
4. Machine Learning: Many algorithms assume Gaussian noise in their models

There's an interesting connection to earlier concepts we discussed: conditional probabilities involving Gaussian distributions have special properties. If you have two variables that follow a joint Gaussian distribution, the conditional distribution of one given the other is also Gaussian - a property that makes these distributions particularly useful in prediction problems.

Chapter 3

Bayesian Filtering Framework

- Bayesian Filtering Framework
 - Recursive state estimation
 - Prediction step (motion model)
 - Update step (measurement model)
 - Chapman-Kolmogorov equation
 - Bayes filter algorithm
 - Linear vs nonlinear systems

3.1 Recursive state estimation

At its heart, recursive state estimation is about continuously updating our belief about a system's state as new measurements come in. Imagine you're trying to track a moving object - at any moment, you want to know its position and velocity, but your sensors are noisy and the object's movement isn't perfectly predictable.

The recursive nature comes from how we update our estimate: instead of processing all past measurements every time we get new data, we maintain a current estimate and update it using only the newest measurement. This makes the process computationally efficient and suitable for real-time applications.

The mathematical framework follows two main steps that repeat over time:

1. Prediction Step (Time Update): First, we predict how the state will evolve based on our system model:

$$p(x_k|z_{1:k-1}) = \int p(x_k|x_{k-1})p(x_{k-1}|z_{1:k-1})dx_{k-1}$$

This integral combines our previous estimate $p(x_{k-1}|z_{1:k-1})$ with our motion model $p(x_k|x_{k-1})$ to predict the new state.

2. Correction Step (Measurement Update): When we get a new measurement, we update our prediction using Bayes' rule:

$$p(x_k|z_{1:k}) = \eta p(z_k|x_k)p(x_k|z_{1:k-1})$$

Here, $p(z_k|x_k)$ is our measurement model, and c is a normalizing constant.

Let's make this concrete with an example. Imagine you're tracking a drone: - State (x): position and velocity - Measurements (z): GPS readings - Motion model: physics equations for how the drone moves - Measurement model: GPS error characteristics

At each time step: 1. You predict where the drone should be based on physics and your last estimate 2. You get a GPS reading 3. You combine your prediction with the measurement to get an updated estimate 4. Repeat for the next time step

The beauty of this approach is that it naturally handles uncertainty. Both your predictions and measurements come with uncertainty (represented as probability distributions), and the Bayesian framework tells you exactly how to combine these uncertainties to get your best estimate.

A key insight is that this framework maintains a complete probability distribution over possible states, not just a single "best guess." This gives you important information about how certain you are about your estimates.

The most famous implementation of this framework is the Kalman Filter, which assumes all uncertainties are Gaussian. This simplifying assumption makes the mathematics tractable while still being useful for many real-world applications. However, the framework itself is more general and can handle non-Gaussian distributions through techniques like particle filters.

3.2 Prediction step (motion model)

The prediction step is fundamentally about using our understanding of how a system evolves over time to estimate its future state. Think of it as using physics to predict where a ball will be in the next moment, given its current position and velocity.

The motion model is mathematically expressed as $p(x(k)|x(k-1))$, which represents the probability of transitioning from state x_{k-1} to state x_k . This probability encapsulates both our deterministic understanding of system dynamics and our uncertainty about random disturbances.

Let's break this down with a concrete example of tracking a moving vehicle in 2D space. Our state vector might include:

- Position (x, y)
- Velocity (vx, vy)
- Acceleration (ax, ay)

In its simplest form, the motion model might use basic physics equations:

$$\begin{aligned}x(k) &= x(k-1) + vx(k-1)\Delta t + \frac{1}{2}ax(k-1)\Delta t^2 \\y(k) &= y(k-1) + vy(k-1)\Delta t + \frac{1}{2}ay(k-1)\Delta t^2 \\vx(k) &= vx(k-1) + ax(k-1)\Delta t \\vy(k) &= vy(k-1) + ay(k-1)\Delta t\end{aligned}$$

However, in real-world scenarios, we need to account for uncertainties. These might come from: 1. Process noise (random disturbances to the system) 2. Model imperfections (our equations aren't perfect) 3. External forces we can't measure directly

This is where the probabilistic nature of the motion model becomes crucial. We typically model these uncertainties using probability distributions. In many cases, we assume Gaussian noise, leading to a linear motion model of the form:

$$x(k) = F(k)x(k-1) + B(k)u(k) + w(k)$$

Where: - $F(k)$ is the state transition matrix (encoding our physics equations) - $B(k)u(k)$ represents known control inputs - $w(k)$ is the process noise (typically assumed Gaussian)

For our vehicle example, the state transition matrix $F(k)$ might look like:

$$\begin{bmatrix} 1 & 0 & \Delta t & 0 & \frac{1}{2}\Delta t^2 & 0 \\ 0 & 1 & 0 & \Delta t & 0 & \frac{1}{2}\Delta t^2 \\ 0 & 0 & 1 & 0 & \Delta t & 0 \\ 0 & 0 & 0 & 1 & 0 & \Delta t \\ 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \end{bmatrix}$$

This matrix encodes how each state variable influences the others over time. Reading across each row tells us how we compute each new state variable.

The complete prediction step involves propagating not just the state estimate but also its uncertainty. If we're using a Gaussian representation, this means we need to update both the mean and covariance of our state estimate:

Mean prediction:

$$\hat{x}(k|k-1) = F(k)\hat{x}(k-1|k-1) + B(k)u(k)$$

Covariance prediction:

$$P(k|k-1) = F(k)P(k-1|k-1)F(k)^\top + Q(k)$$

Where $Q(k)$ is the process noise covariance matrix that quantifies our uncertainty about the motion model.

This probabilistic approach allows us to: 1. Make predictions about future states 2. Maintain an estimate of our uncertainty 3. Account for both systematic and random effects 4. Prepare for the measurement update step where we'll combine these predictions with actual measurements

3.3 Update step (measurement model)

The measurement update (or correction step) is where we refine our predicted state estimate by incorporating new sensor measurements. This step is fundamentally based on Bayes' rule, which tells us how to update probabilities when we get new evidence.

The key equation for the measurement update is:

$$p(x_k|z_{1:k}) = \eta p(z_k|x_k)p(x_k|z_{1:k-1})$$

Let's break this down using our vehicle tracking example. Imagine we have GPS measurements that give us position information. The measurement model $p(z_k|x_k)$ describes how our sensor readings relate to the true state, including sensor noise and limitations.

In its simplest form, for a linear system with Gaussian noise, the measurement model can be written as:

$$z(k) = H(k)x(k) + v(k)$$

Where: - $z(k)$ is the measurement vector - $H(k)$ is the measurement matrix that maps the state space to measurement space - $v(k)$ is the measurement noise (typically assumed Gaussian)

For our vehicle tracking example with GPS measurements, the measurement matrix $H(k)$ might look like:

$$\begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 \end{bmatrix}$$

This matrix indicates that we're only measuring position (x, y) , not velocity or acceleration.

The actual update computation involves several steps:

1. First, we compute the innovation (the difference between predicted and actual measurements):

$$y(k) = z(k) - H(k)\hat{x}(k|k-1)$$

2. Then, we calculate the innovation covariance:

$$S(k) = H(k)P(k|k-1)H(k)^\top + R(k)$$

where $R(k)$ is the measurement noise covariance matrix

3. Next, we compute the Kalman gain, which determines how much we trust our measurement versus our prediction:

$$K(k) = P(k|k-1)H(k)^\top S(k)^{-1}$$

4. Finally, we update our state estimate and its covariance:

$$\begin{aligned}\hat{x}(k|k) &= \hat{x}(k|k-1) + K(k)y(k) \\ P(k|k) &= (I - K(k)H(k))P(k|k-1)\end{aligned}$$

The Kalman gain $K(k)$ is particularly interesting because it acts as a weighting factor. When our measurements are very precise (small R), the Kalman gain will be

larger, meaning we trust the measurements more. When our measurements are noisy (large R), the gain will be smaller, and we'll trust our predictions more.

Consider what happens in extreme cases: - If our GPS suddenly becomes very accurate ($R \rightarrow 0$), K will increase, and we'll trust the GPS more - If our GPS is experiencing interference (large R), K will decrease, and we'll rely more on our motion model

This adaptive behavior is what makes recursive state estimation so powerful. The system automatically balances between prediction and measurement based on their relative uncertainties.

An important practical consideration is choosing appropriate values for the measurement noise covariance $R(k)$. This matrix represents our understanding of sensor characteristics and limitations. For a GPS sensor, we might set:

$$\begin{bmatrix} \sigma_x^2 & 0 \\ 0 & \sigma_y^2 \end{bmatrix}$$

where σ_x and σ_y represent our uncertainty in x and y measurements.

3.4 Chapman-Kolmogorov equation

The Chapman-Kolmogorov equation is a fundamental concept in probability theory and plays a crucial role in state estimation. This equation describes how probability distributions evolve over time in a Markov process.

The Chapman-Kolmogorov equation is mathematically expressed as:

$$p(x_{k+1}|z_{1:k}) = \int p(x_{k+1}|x_k)p(x_k|z_{1:k})dx_k$$

Think of this equation as describing how we can “step forward” in time with our probability distributions. Let's break down what each part means and why it's important.

The left side, $p(x_{k+1}|z_{1:k})$, represents our prediction of the state at time $k+1$, given all measurements up to time k . This is what we want to calculate.

On the right side, we have two key components:

1. $p(x_{k+1}|x_k)$ is our motion model - how the state evolves from one time step to the next
2. $p(x_k|z_{1:k})$ is our current belief about the state at time k

The integral combines these components across all possible current states. It's like considering every possible current state, figuring out where it might lead, and weighting those possibilities by how likely we think each current state is.

Let's make this concrete with an example. Imagine tracking a car on a one-dimensional road:

- Current position is x
- Future position is x

- We have some uncertainty about both

The Chapman-Kolmogorov equation tells us to:

1. Consider each possible current position
2. For each current position, consider all possible future positions
3. Weight each possibility by how likely we think it is
4. Sum up all these weighted possibilities

This process naturally handles uncertainty propagation. If we're very uncertain about the current state, this uncertainty will be reflected in our prediction through the integration process.

The equation becomes particularly elegant when working with Gaussian distributions. In this case, if:

- Current belief is Gaussian with mean μ_k and variance σ_k^2
- Motion model is Gaussian with mean shift δ and variance τ^2

Then the prediction will also be Gaussian with:

- Mean: $\mu_{k+1} = \mu_k + \delta$
- Variance: $\sigma_{k+1}^2 = \sigma_k^2 + \tau^2$

This shows how uncertainties add up as we make predictions further into the future.

The Chapman-Kolmogorov equation is particularly important because:

1. It forms the theoretical foundation for the prediction step in Bayesian filtering
2. It respects the Markov property we discussed earlier
3. It provides a mathematically rigorous way to propagate uncertainties
4. It connects continuous and discrete-time processes

In practical applications, we often can't solve the integral analytically. This leads to various approximation methods:

- Kalman filters use Gaussian approximations
- Particle filters use numerical sampling
- Grid-based methods discretize the state space

3.5 Bayes filter algorithm

The Bayes filter is a probabilistic approach to estimating the state of a dynamic system over time using noisy measurements. Think of it as a mathematical framework for maintaining an educated guess about what's happening in a system, constantly refining that guess as new information arrives.

The core principle rests on maintaining a belief state - a probability distribution over all possible states. This belief state represents our uncertainty about the true state of the system. Let's break down how this works through the two main steps that occur recursively:

The Prediction Step (Time Update): In this first phase, we predict how our system will evolve based on our understanding of its dynamics. Imagine tracking a flying drone - even without looking at it, we can predict where it should be based on physics and our last known information about its position and velocity. This step uses the Chapman-Kolmogorov equation we discussed earlier:

$$p(x_k|z_{1:k-1}) = \int p(x_k|x_{k-1})p(x_{k-1}|z_{1:k-1})dx_{k-1}$$

This equation tells us to consider all possible previous states and how they might evolve into current states. The result is a prediction that accounts for all uncertainties in the system's dynamics.

The Correction Step (Measurement Update): When we get new sensor information, we update our prediction using Bayes' rule:

$$p(x_k|z_{1:k}) = \eta p(z_k|x_k)p(x_k|z_{1:k-1})$$

Here, we're weighing our prediction against new evidence, much like a detective updating their theory based on new clues.

Let's make this concrete with an example of a self-driving car:

Starting State:

- The car believes it's at a certain position with some uncertainty
- This belief is represented as a probability distribution over possible positions

Prediction Phase:

- The car knows it's moving forward at 30 mph
- Using this information and basic physics, it predicts where it should be after a small time interval
- The uncertainty grows during this prediction (the car might be sliding slightly, or its speedometer might be imperfect)

Measurement Phase:

- The car's GPS provides a new position reading
- The car's cameras detect lane markers
- These measurements are combined with the prediction to form an updated belief
- The uncertainty typically decreases during this phase as new evidence arrives

The beauty of the Bayes filter lies in how it handles uncertainty:

- If sensors are very accurate, their measurements are weighted more heavily
- If the motion is very predictable, the predictions are trusted more
- The algorithm automatically balances these factors based on their relative uncertainties

Real-world implementations often make specific assumptions about the nature of uncertainties and system dynamics. The most famous variant is the Kalman filter, which assumes:

- Linear system dynamics
- Gaussian uncertainties
- Additive noise

However, the general Bayes filter framework can handle non-linear systems and non-Gaussian uncertainties through variants like:

- Extended Kalman Filter: Handles mild non-linearities through linearization
- Unscented Kalman Filter: Better handles non-linear systems
- Particle Filter: Can handle any type of uncertainty or dynamics

3.6 Linear vs nonlinear systems

In the context of state estimation, a system's linearity or nonlinearity affects how states evolve over time and how measurements relate to states. This distinction is crucial because it determines which filtering approaches we can use effectively.

3.6.1 Linear Systems

A system is linear if it satisfies two key properties: superposition and homogeneity. In state estimation terms, this means:

For the motion model, a linear system follows the form:

$$x(k+1) = Ax(k) + Bu(k) + w(k)$$

Where:

- A is the state transition matrix
- B is the control input matrix
- $w(k)$ is process noise
- All relationships between variables are strictly linear

For the measurement model, linearity means:

$$z(k) = Hx(k) + v(k)$$

Where:

- H is the measurement matrix
- $v(k)$ is measurement noise

Consider tracking a train moving along a straight track at constant speed. This is approximately linear because:

- Position changes linearly with time
- Velocity remains constant
- Measurements (like position from track sensors) are directly proportional to state

3.6.2 Nonlinear Systems

Real-world systems are often nonlinear. The state evolution or measurements might involve:

- Trigonometric functions
- Quadratic terms
- Products of state variables
- Any other nonlinear mathematical relationships

The general form becomes:

$$\begin{aligned}x(k+1) &= f(x(k), u(k)) + w(k) \\ z(k) &= h(x(k)) + v(k)\end{aligned}$$

Where $f()$ and $h()$ are nonlinear functions.

Consider tracking an aircraft:

- Position updates involve trigonometric functions of orientation
- Aerodynamic forces are quadratic with velocity
- Radar measurements give range and bearing, requiring nonlinear conversions to Cartesian coordinates

The implications for filtering are profound:

For Linear Systems:

- The Kalman Filter provides an optimal solution
- Uncertainties remain Gaussian if noise is Gaussian
- Computations are relatively simple and fast
- Results are guaranteed to converge under certain conditions

For Nonlinear Systems:

- The basic Kalman Filter no longer works optimally
- Uncertainties may become non-Gaussian even with Gaussian noise
- We need more sophisticated approaches:
 - Extended Kalman Filter (EKF): Linearizes around current estimate
 - Unscented Kalman Filter (UKF): Uses carefully chosen sample points
 - Particle Filter: Represents uncertainty with discrete particles

Let's consider a specific example: a pendulum.

- Linear approximation works when swing angle is small ($\sin(\theta) \approx \theta$)
- As swing amplitude increases, nonlinear effects become important
- The true motion involves trigonometric functions of angle

This demonstrates how real systems might be approximated as linear within certain operating ranges but require nonlinear treatment for full accuracy.

Chapter 4

Kalman Filtering

- Kalman Filtering
 - Linear Kalman Filter
 - System model and assumptions
 - Prediction equations
 - Update equations
 - Uncertainty propagation
 - Extended Kalman Filter (EKF)
 - Linearization process
 - Jacobian matrices
 - Algorithm implementation
 - Unscented Kalman Filter (UKF)
 - Sigma points
 - Unscented transform
 - Algorithm implementation

4.1 Linear Kalman Filter

4.1.1 System model and assumptions

At its core, the linear Kalman filter describes a system using two fundamental equations:

1. The State Equation (also called the Process Model):

$$x(k) = F(k)x(k-1) + B(k)u(k) + w(k)$$

This equation tells us how the system evolves over time. Here, $x(k)$ is the current state, $F(k)$ is the state transition matrix that describes how the state naturally evolves, $B(k)u(k)$ represents known control inputs, and $w(k)$ is the process noise.

2. The Measurement Equation (also called the Observation Model):

$$z(k) = H(k)x(k) + v(k)$$

This equation describes how we observe the system. $z(k)$ represents our measurements, $H(k)$ is the measurement matrix that maps the state to measurements, and $v(k)$ is the measurement noise.

Now, let's examine the crucial assumptions that make the linear Kalman filter work:

First, we assume linearity in both equations. This means that both the state transitions and measurements must be linear functions of the state. In real-world terms, if you double the state, the output should double too. This is why it's called a "linear" Kalman filter.

Second, we make important assumptions about the noise terms $w(k)$ and $v(k)$. We assume they are:

- Zero-mean Gaussian white noise
- Uncorrelated with each other
- Uncorrelated in time (white noise)
- Known covariances (Q for process noise and R for measurement noise)

Think of these noise assumptions like rolling fair dice repeatedly - each roll is independent, and while you might not know the exact outcome, you know the probability distribution.

Third, we assume we have some initial knowledge about the state, typically expressed as an initial state estimate $x(0)$ and its uncertainty $P(0)$. This doesn't need to be perfectly accurate, but better initial estimates generally lead to faster convergence.

Fourth, we assume that the matrices $F(k)$, $H(k)$, $B(k)$, $Q(k)$, and $R(k)$ are known at each time step. They can vary with time, but we need to know their values.

The beauty of these assumptions is that they allow us to derive optimal state estimates in a mathematically rigorous way. When these assumptions hold true, the Kalman filter gives us the best possible estimate of the state (in terms of minimum mean square error).

4.1.2 Prediction equations

Let's go through the prediction equations of the Kalman filter. We'll break this down step by step, with a focus on what each equation means practically.

The Prediction Phase (Time Update): Let's start with predicting where our system will be in the next time step. We have two key equations here:

1. State Prediction:

$$x^-(k) = F(k)x(k-1) + B(k)u(k)$$

Where:

- $x^-(k)$ is our predicted state (the superscript minus means "before measurement update")
- $x(k-1)$ is our previous state estimate
- $F(k)$ is the state transition matrix
- $B(k)$ is the control input matrix
- $u(k)$ is the control input vector

Think of this like predicting where a moving object will be based on its current position, velocity, and any forces we're applying to it.

2. Error Covariance Prediction:

$$P^-(k) = F(k)P(k-1)F(k)^T + Q(k)$$

Where:

- $P^-(k)$ is our predicted error covariance
- $P(k-1)$ is our previous error covariance
- $F(k)^T$ is the transpose of $F(k)$
- $Q(k)$ is the process noise covariance

This equation tells us how uncertain we are about our prediction. It's crucial for engineering applications because it helps us understand the reliability of our estimates.

Practical Engineering Considerations:

1. State Transition Matrix (F):

For a simple position-velocity system in one dimension, F might look like:

1	$F = \begin{bmatrix} 1 & \Delta t \end{bmatrix}$
2	$\begin{bmatrix} 0 & 1 \end{bmatrix}$

Where Δt is your sampling time. This models constant velocity motion.

2. Process Noise (Q):

For the same system, a common model is:

1	$Q = \begin{bmatrix} \Delta[(\Delta t/4)^2] & \Delta[(\Delta t^3/2)^2] \end{bmatrix}$
2	$\begin{bmatrix} \Delta[(\Delta t^3/2)^2] & \Delta[\Delta t^2] \end{bmatrix}$

Where σ^2 is the variance of your acceleration noise.

3. Control Input:

If you're applying known forces or controls, $B(k)u(k)$ accounts for these. For example, in a rocket system, this might represent thrust.

Engineering Tips:

1. Always check units! Make sure your matrices are dimensionally consistent.
2. When implementing, start with small time steps (Δt) to reduce linearization errors.
3. Q should be tuned based on your system's characteristics - if you're unsure, start with a diagonal matrix with conservative (larger) values.

Common Pitfalls to Avoid:

1. Don't forget to propagate uncertainties - P is just as important as \hat{x}
2. Watch out for numerical issues - consider using square root forms for P if precision is critical
3. Make sure F and Q are properly synchronized with your sampling time

4.1.3 Update equations

The Measurement Update (Correction Phase) consists of three key equations that work together to incorporate new measurement information into our state estimate:

1. The Kalman Gain Equation:

$$K(k) = P^-(k)H(k)^T[H(k)P^-(k)H(k)^T + R(k)]^{-1}$$

This equation determines how much we trust our new measurement versus our prediction. Think of it as a weighting factor that balances between our model and our sensors. Let's break down what each term means:

- $P^-(k)$ is our predicted error covariance (from the prediction step)
- $H(k)$ is our measurement matrix
- $R(k)$ is our measurement noise covariance
- The $^{-1}$ indicates matrix inversion

2. The State Update Equation:

$$x(k) = x^-(k) + K(k)[z(k) - H(k)x^-(k)]$$

This equation updates our state estimate based on the measurement. The term $[z(k) - H(k)\hat{x}^-(k)]$ is called the innovation or measurement residual - it represents the difference between what we measured and what we expected to measure. We then weight this difference by the Kalman gain to determine how much to correct our prediction.

3. The Error Covariance Update:

$$P(k) = [I - K(k)H(k)]P^-(k)$$

This equation updates our uncertainty estimate. The term $[I - K(k)H(k)]$ is sometimes called the stability factor, as it helps ensure numerical stability in our computations.

From an engineering perspective, here are some crucial insights about these equations:

The Kalman gain $K(k)$ will automatically adapt based on the relative uncertainties. If $R(k)$ is large (noisy measurements), $K(k)$ will be smaller, meaning we trust our predictions more. If $P^-(k)$ is large (uncertain predictions), $K(k)$ will be larger, meaning we trust the measurements more.

For implementation, you might encounter alternate forms of the covariance update equation:

- Joseph form:

$$P(k) = [I - K(k)H(k)]P^-(k)[I - K(k)H(k)]^T + K(k)R(k)K(k)^T$$

This form is more computationally expensive but numerically more stable. It's particularly useful when dealing with ill-conditioned matrices.

A practical example: Consider a simple position tracking system where we measure position directly. Our measurement matrix might be:

```
1 H = [1 0]
```

This indicates we're measuring position but not velocity directly. Our measurement noise covariance R might be a single value representing our sensor's variance:

```
1 R = [_sensor]
```

4.1.4 Uncertainty propagation

At its core, uncertainty propagation in Kalman filtering deals with how our uncertainty about the state evolves over time and through measurements. This uncertainty is represented by the covariance matrix P, which is a symmetric matrix where the diagonal elements represent variances of individual state components, and off-diagonal elements represent their correlations.

Let's start with linear uncertainty propagation through the prediction step:

$$P^-(k) = F(k)P(k-1)F(k)^T + Q(k)$$

This equation comes from the linear transformation of random variables. When we have a random variable x and transform it linearly by A, the covariance transforms as APA. In our case, F is our transformation matrix, and we add Q to account for additional uncertainty from process noise.

To understand this deeper, let's consider a simple 2D state space of position and velocity:

```
1 P = [_pos      _pos_vel ]
2     [_pos_vel  ^2_vel   ]
```

When we propagate this through time with:

```
1 F = [1  Δt]
2     [0  1]
```

The uncertainty grows in a specific pattern. The position uncertainty increases due to both:

1. The existing position uncertainty
2. The velocity uncertainty (scaled by Δt)
3. The correlation between position and velocity
4. The process noise Q

This creates a characteristic “banana-shaped” uncertainty region in position-velocity space, because position and velocity become correlated through the prediction step even if they weren't initially.

Moving to the measurement update, uncertainty reduction happens through:

$$P(k) = [I - K(k)H(k)]P^-(k)$$

This equation shows how incorporating a measurement reduces our uncertainty. The amount of reduction depends on:

1. How uncertain we were before (P)
2. How accurate our measurement is (R)
3. What we're measuring (H)

The Kalman gain K plays a crucial role here. It's derived to minimize the trace of $P(k)$, which means it minimizes the sum of variances of our state estimates. This is why the Kalman filter is called an optimal estimator - it provides the minimum variance estimate given our assumptions.

A practical example helps illustrate this: Imagine tracking a moving object where:

- Position measurements have uncertainty $\sigma_{\text{meas}} = 1\text{m}$
- Initial velocity uncertainty $\sigma_{\text{vel}} = 0.1\text{ m/s}$
- Sampling time $\Delta t = 0.1\text{s}$

Our prediction step would propagate uncertainty like this:

```
1 # Initial uncertainty
2 P = [[1.0, 0.0],      # Position variance = 1m²
3       [0.0, 0.01]]   # Velocity variance = 0.01(m/s)²
4
5 # State transition
6 F = [[1.0, 0.1],     # Δt = 0.1s
7       [0.0, 1.0]]
8
9 # Process noise (simplified)
10 Q = [[0.001, 0.0],  # Small position process noise
11       [0.0, 0.001]] # Small velocity process noise
12
13 # Predicted uncertainty
14 P_pred = F @ P @ F.T + Q
```

After prediction, you'll notice:

1. Position uncertainty has increased
2. A correlation has developed between position and velocity
3. Velocity uncertainty has grown slightly due to Q

This understanding of uncertainty propagation is crucial for:

1. Filter tuning - setting appropriate Q and R values
2. Sensor fusion - knowing how to weight different measurements
3. System design - understanding how measurement frequency affects estimation quality

4.2 Extended Kalman Filter (EKF)

4.2.1 Linearization process

Through the Extended Kalman Filter (EKF) linearization process, we adapt the linear Kalman filter concepts to handle nonlinear systems.

In the EKF, we deal with nonlinear system equations:

State Evolution (Process Model):

$$x(k) = f(x(k-1), u(k)) + w(k)$$

Measurement Model:

$$z(k) = h(x(k)) + v(k)$$

Where $f()$ and $h()$ are nonlinear functions. The key insight of the EKF is that we can approximate these nonlinear functions using Taylor series expansion around our current estimate. Let's break this down step by step:

Linearization Process

We need to compute Jacobian matrices - these are matrices of partial derivatives that give us the best linear approximation of our nonlinear functions at a specific point. We compute:

$$F(k) = \partial f / \partial x \text{ evaluated at } x(k-1) \quad H(k) = \partial h / \partial x \text{ evaluated at } x^-(k)$$

Let's work through a practical example. Consider a robot moving in 2D with state vector:

$$x = [x\text{position}, y\text{position}, \text{headingangle}\theta]$$

The nonlinear process model might be:

```
1 def f(x, u):
2     # x: state [x, y, ]
3     # u: control [velocity, angular_velocity]
4     dt = 0.1 # time step
5
6     x_new = x[0] + u[0]*cos(x[2])*dt
7     y_new = x[1] + u[0]*sin(x[2])*dt
8     _new = x[2] + u[1]*dt
9
10    return np.array([x_new, y_new, _new])
```

To linearize this, we compute the Jacobian F:

```
1 def compute_F(x, u):
2     dt = 0.1
3     v = u[0] # velocity
4
5     F = np.array([
6         [1, 0, -v*sin(x[2])*dt],
7         [0, 1, v*cos(x[2])*dt],
8         [0, 0, 1]
9     ])
10    return F
```


Using the Linearized Models

Once we have our Jacobians, the EKF prediction equations become:

State Prediction:

$$x^-(k) = f(x(k-1), u(k)) // U senonlinear function$$

Covariance Prediction:

$$P^-(k) = F(k)P(k-1)F(k)^T + Q(k) // U selinearized F$$

The measurement update is similar:

Innovation:

$$y(k) = z(k) - h(x^-(k)) // U senonlinear function$$

Kalman Gain:

$$K(k) = P^-(k)H(k)^T[H(k)P^-(k)H(k)^T + R(k)]^{-1} // U selinearized H$$

State Update:

$$x(k) = x^-(k) + K(k)y(k)$$

Covariance Update:

$$P(k) = [I - K(k)H(k)]P^-(k)$$

Important Considerations

- The linearization is only valid near the operating point. If your estimate strays too far from the true state, the approximation breaks down.
- You need to recompute the Jacobians at each time step since they depend on the current state estimate.
- Numerical computation of Jacobians can be useful for complex systems:

```
1 def numerical_jacobian(f, x, dx=1e-6):
2     n = len(x)
3     J = np.zeros((n, n))
4     for i in range(n):
5         x_plus = x.copy()
6         x_plus[i] += dx
7         x_minus = x.copy()
8         x_minus[i] -= dx
9         J[:, i] = (f(x_plus) - f(x_minus)) / (2*dx)
10    return J
```

4.2.2 Jacobian matrices

The Jacobian matrices in EKF serve as our linear approximation of nonlinear functions. We need two key Jacobians: $F(k)$ for the process model and $H(k)$ for the measurement model.

The Process Model Jacobian (F)

$F(k)$ represents how small changes in our state affect the predicted next state. Mathematically, it's the matrix of partial derivatives of our process function f with respect to each state variable:

$$F(k) = \partial f / \partial x = [\partial f_1 / \partial x_1 \partial f_1 / \partial x_2 \partial f_1 / \partial x_3 \dots] [\partial f_2 / \partial x_1 \partial f_2 / \partial x_2 \partial f_2 / \partial x_3 \dots] [\partial f_3 / \partial x_1 \partial f_3 / \partial x_2 \partial f_3 / \partial x_3 \dots]$$

Let's work through a concrete example. Consider a robot moving in 2D with state vector $x = [x, y, \theta]$, where:

- x, y are position coordinates
- θ is the heading angle
- The robot moves with velocity v and angular velocity ω

The nonlinear process model would be:

$$x(k+1) = x(k) + v * \cos(\theta) * dt \quad y(k+1) = y(k) + v * \sin(\theta) * dt \quad \theta(k+1) = \theta(k) + \omega * dt$$

The Jacobian F would then be:

```
1 def compute_process_jacobian(x, v, dt):
2     """
3     Compute the process model Jacobian for a 2D robot
4     x: State vector [x, y, ]
5     v: Linear velocity
6     dt: Time step
7     """
8     F = np.array([
9         [1, 0, -v*sin() * dt], # x' / x, x' / y, x' /
10        [0, 1, v*cos() * dt],  # y' / x, y' / y, y' /
11        [0, 0, 1]               #   ' / x,   ' / y,   ' /
12    ])
13    return F
```

The Measurement Model Jacobian (H)

$H(k)$ represents how small changes in our state affect what we expect to measure. It's the matrix of partial derivatives of our measurement function h with respect to each state variable:

$$H(k) = \partial h / \partial x = [\partial h_1 / \partial x_1 \partial h_1 / \partial x_2 \partial h_1 / \partial x_3 \dots] [\partial h_2 / \partial x_1 \partial h_2 / \partial x_2 \partial h_2 / \partial x_3 \dots] [\dots \dots \dots]$$

For example, if we're measuring range and bearing to a landmark at position $[x_l, y_l]$:

$$r = \sqrt{(x - x_l)^2 + (y - y_l)^2}$$

$$\beta = \text{atan2}(y_l - y, x_l - x) - \theta$$

The measurement Jacobian would be:

```

1 def compute_measurement_jacobian(x, landmark_pos):
2     """
3     Compute measurement Jacobian for range-bearing measurements
4     x: State vector [x, y, ]
5     landmark_pos: Position of landmark [x, y]
6     """
7     dx = landmark_pos[0] - x[0]
8     dy = landmark_pos[1] - x[1]
9     r = np.sqrt(dx**2 + dy**2)
10
11     H = np.array([
12         [-dx/r,      -dy/r,      0      ], # r / x,   r / y,   r /
13         [dy/r**2,     -dx/r**2,    -1     ] #   / x,   / y,   /
14     ])
15     return H

```

Important Considerations for Jacobian Computation:

1. Singularity Points: Some configurations might lead to undefined Jacobians. For example, in our range-bearing case, when the robot is exactly at the landmark position ($r = 0$), the bearing Jacobian becomes undefined.
2. Numerical Stability: When implementing Jacobians, watch out for numerical issues. For very small values, you might need to add small constants to denominators:

```

1 def safe_division(num, den, eps=1e-10):
2     return num / (den + eps)

```

3. Verification: You can verify your analytical Jacobians using numerical differentiation:

```

1 def verify_jacobian(f, x, dx=1e-6):
2     """
3     Verify analytical Jacobian against numerical approximation
4     """
5     analytical = compute_jacobian(x)
6     numerical = numerical_jacobian(f, x, dx)
7     difference = np.abs(analytical - numerical)
8     print("Maximum difference:", np.max(difference))

```

4.2.3 Algorithm implementation

Let's create a comprehensive implementation of both Kalman Filter and Extended Kalman Filter with detailed documentation and test scenarios. The key components of this implementation are:

Base Classes

- `KalmanFilterState`: A dataclass to store filter states
- `KalmanFilter`: Implementation of the linear Kalman filter
- `ExtendedKalmanFilter`: Implementation of the EKF

Test Scenarios

- Linear case: Tracking an object moving with constant velocity
- Nonlinear case: Tracking a robot with nonlinear dynamics using range-bearing measurements

The code includes detailed comments explaining each component and step.

```
1 import numpy as np
2 from dataclasses import dataclass
3 from typing import Optional, Tuple, Callable
4 import matplotlib.pyplot as plt
5
6 @dataclass
7 class KalmanFilterState:
8     """
9     Dataclass to store the state of a Kalman Filter
10
11     Attributes:
12         x (np.ndarray): State estimate vector
13         P (np.ndarray): State covariance matrix
14         dim_x (int): Dimension of state vector
15         dim_z (int): Dimension of measurement vector
16     """
17     x: np.ndarray # State estimate
18     P: np.ndarray # Covariance matrix
19     dim_x: int # State dimension
20     dim_z: int # Measurement dimension
21
22 class KalmanFilter:
23     """
24     Implementation of a linear Kalman Filter.
25
26     This implementation follows the standard Kalman Filter equations for
27     linear systems with Gaussian noise.
28     """
29
30     def __init__(self, dim_x: int, dim_z: int):
31         """
32         Initialize Kalman Filter
33
34         Args:
35             dim_x: Dimension of state vector
36             dim_z: Dimension of measurement vector
37         """
38         self.dim_x = dim_x
39         self.dim_z = dim_z
40
41         # Initialize matrices
42         self.x = np.zeros(dim_x) # State estimate
43         self.P = np.eye(dim_x) # State covariance
44         self.F = np.eye(dim_x) # State transition matrix
```

```

45     self.H = np.zeros((dim_z, dim_x)) # Measurement matrix
46     self.R = np.eye(dim_z)           # Measurement noise covariance
47     self.Q = np.eye(dim_x)           # Process noise covariance
48
49     def predict(self, u: Optional[np.ndarray] = None, B: Optional[np.
ndarray] = None) -> None:
50         """
51         Predict step of the Kalman Filter
52
53         Args:
54             u: Optional control input
55             B: Optional control matrix
56         """
57         # State prediction
58         if u is not None and B is not None:
59             self.x = self.F @ self.x + B @ u
60         else:
61             self.x = self.F @ self.x
62
63         # Covariance prediction
64         self.P = self.F @ self.P @ self.F.T + self.Q
65
66     def update(self, z: np.ndarray) -> None:
67         """
68         Update step of the Kalman Filter
69
70         Args:
71             z: Measurement vector
72         """
73         # Innovation and innovation covariance
74         y = z - self.H @ self.x
75         S = self.H @ self.P @ self.H.T + self.R
76
77         # Kalman gain
78         K = self.P @ self.H.T @ np.linalg.inv(S)
79
80         # State and covariance update
81         self.x = self.x + K @ y
82         self.P = (np.eye(self.dim_x) - K @ self.H) @ self.P
83
84     def get_state(self) -> KalmanFilterState:
85         """Return current filter state"""
86         return KalmanFilterState(
87             x=self.x.copy(),
88             P=self.P.copy(),
89             dim_x=self.dim_x,
90             dim_z=self.dim_z
91         )
92
93     class ExtendedKalmanFilter:
94         """
95         Implementation of an Extended Kalman Filter.
96
97         This implementation handles nonlinear systems by linearizing
98         around the current state estimate.
99         """
100
101     def __init__(self, dim_x: int, dim_z: int,

```

```

102         f: Callable, h: Callable,
103         compute_F: Callable, compute_H: Callable):
104     """
105     Initialize Extended Kalman Filter
106
107     Args:
108         dim_x: Dimension of state vector
109         dim_z: Dimension of measurement vector
110         f: Nonlinear state transition function
111         h: Nonlinear measurement function
112         compute_F: Function to compute state transition Jacobian
113         compute_H: Function to compute measurement Jacobian
114     """
115     self.dim_x = dim_x
116     self.dim_z = dim_z
117
118     # Nonlinear functions
119     self.f = f          # State transition function
120     self.h = h          # Measurement function
121     self.compute_F = compute_F # State transition Jacobian
122     self.compute_H = compute_H # Measurement Jacobian
123
124     # Initialize matrices
125     self.x = np.zeros(dim_x)
126     self.P = np.eye(dim_x)
127     self.R = np.eye(dim_z)
128     self.Q = np.eye(dim_x)
129
130     def predict(self, u: Optional[np.ndarray] = None) -> None:
131         """
132         Predict step of the EKF
133
134         Args:
135             u: Optional control input
136         """
137         # Compute state transition Jacobian
138         F = self.compute_F(self.x, u)
139
140         # Nonlinear state prediction
141         if u is not None:
142             self.x = self.f(self.x, u)
143         else:
144             self.x = self.f(self.x, None)
145
146         # Covariance prediction using linearized F
147         self.P = F @ self.P @ F.T + self.Q
148
149     def update(self, z: np.ndarray) -> None:
150         """
151         Update step of the EKF
152
153         Args:
154             z: Measurement vector
155         """
156         # Compute measurement Jacobian
157         H = self.compute_H(self.x)
158
159         # Innovation using nonlinear measurement function

```

```

160     y = z - self.h(self.x)
161
162     # Innovation covariance using linearized H
163     S = H @ self.P @ H.T + self.R
164
165     # Kalman gain
166     K = self.P @ H.T @ np.linalg.inv(S)
167
168     # State and covariance update
169     self.x = self.x + K @ y
170     self.P = (np.eye(self.dim_x) - K @ H) @ self.P
171
172     def get_state(self) -> KalmanFilterState:
173         """Return current filter state"""
174         return KalmanFilterState(
175             x=self.x.copy(),
176             P=self.P.copy(),
177             dim_x=self.dim_x,
178             dim_z=self.dim_z
179         )
180
181     # Test scenarios
182
183     def test_linear_constant_velocity():
184         """
185         Test scenario: Track an object moving with constant velocity
186         using a linear Kalman filter
187         """
188         # Initialize Kalman Filter
189         kf = KalmanFilter(dim_x=4, dim_z=2) # State: [x, y, vx, vy],
190         Measurement: [x, y]
191         dt = 0.1
192
193         # Set up system matrices
194         kf.F = np.array([
195             [1, 0, dt, 0],
196             [0, 1, 0, dt],
197             [0, 0, 1, 0],
198             [0, 0, 0, 1]
199         ])
200
201         kf.H = np.array([
202             [1, 0, 0, 0],
203             [0, 1, 0, 0]
204         ])
205
206         # Set noise covariances
207         kf.R = np.eye(2) * 0.1 # Measurement noise
208         kf.Q = np.eye(4) * 0.01 # Process noise
209
210         # Generate true trajectory
211         t = np.arange(0, 10, dt)
212         true_x = 0.1 * t**2 # Quadratic motion in x
213         true_y = 0.1 * t     # Linear motion in y
214
215         # Generate noisy measurements
216         measurements = np.column_stack((
217             true_x + np.random.normal(0, 0.1, len(t)),

```

```

217         true_y + np.random.normal(0, 0.1, len(t))
218     ))
219
220     # Run filter
221     estimated_states = []
222     for z in measurements:
223         kf.predict()
224         kf.update(z)
225         estimated_states.append(kf.get_state().x.copy())
226
227     estimated_states = np.array(estimated_states)
228
229     # Plot results
230     plt.figure(figsize=(12, 6))
231
232     plt.subplot(121)
233     plt.plot(true_x, true_y, 'k-', label='True')
234     plt.plot(measurements[:, 0], measurements[:, 1], 'r.', label='
Measurements')
235     plt.plot(estimated_states[:, 0], estimated_states[:, 1], 'b--', label='
Estimated')
236     plt.legend()
237     plt.title('Position Track')
238     plt.xlabel('X')
239     plt.ylabel('Y')
240
241     plt.subplot(122)
242     plt.plot(t, estimated_states[:, 2], 'r-', label='Estimated Vx')
243     plt.plot(t, estimated_states[:, 3], 'b-', label='Estimated Vy')
244     plt.legend()
245     plt.title('Velocity Estimates')
246     plt.xlabel('Time')
247     plt.ylabel('Velocity')
248
249     plt.tight_layout()
250     plt.show()
251
252 def test_nonlinear_robot():
253     """
254     Test scenario: Track a robot moving with nonlinear dynamics
255     using an Extended Kalman Filter
256     """
257     def f(x, u):
258         """Nonlinear state transition function"""
259         dt = 0.1
260         if u is None:
261             v = 1.0 # Constant velocity if no control
262             omega = 0.1 # Constant angular velocity
263         else:
264             v = u[0]
265             omega = u[1]
266
267         theta = x[2]
268         return np.array([
269             x[0] + v * np.cos(theta) * dt,
270             x[1] + v * np.sin(theta) * dt,
271             x[2] + omega * dt
272         ])

```



```

273
274 def h(x):
275     """Nonlinear measurement function (range and bearing to landmark)"""
276     "
277     landmark = np.array([5.0, 5.0]) # Fixed landmark position
278     dx = landmark[0] - x[0]
279     dy = landmark[1] - x[1]
280
281     r = np.sqrt(dx**2 + dy**2)
282     bearing = np.arctan2(dy, dx) - x[2]
283     return np.array([r, bearing])
284
285 def compute_F(x, u):
286     """Compute state transition Jacobian"""
287     dt = 0.1
288     if u is None:
289         v = 1.0
290     else:
291         v = u[0]
292
293     theta = x[2]
294     return np.array([
295         1, 0, -v * np.sin(theta) * dt],
296         [0, 1, v * np.cos(theta) * dt],
297         [0, 0, 1]
298     ])
299
300 def compute_H(x):
301     """Compute measurement Jacobian"""
302     landmark = np.array([5.0, 5.0])
303     dx = landmark[0] - x[0]
304     dy = landmark[1] - x[1]
305     r = np.sqrt(dx**2 + dy**2)
306
307     return np.array([
308         -dx/r, -dy/r, 0],
309         [dy/r**2, -dx/r**2, -1]
310     ])
311
312 # Initialize EKF
313 ekf = ExtendedKalmanFilter(
314     dim_x=3, # State: [x, y, theta]
315     dim_z=2, # Measurement: [range, bearing]
316     f=f,
317     h=h,
318     compute_F=compute_F,
319     compute_H=compute_H
320 )
321
322 # Set noise covariances
323 ekf.R = np.diag([0.1, 0.1]) # Measurement noise
324 ekf.Q = np.diag([0.01, 0.01, 0.01]) # Process noise
325
326 # Generate true trajectory
327 t = np.arange(0, 10, 0.1)
328 true_states = []
329 x = np.array([0.0, 0.0, 0.0]) # Initial state

```

```

330     for _ in t:
331         true_states.append(x.copy())
332         x = f(x, None) # Use constant velocity and angular velocity
333
334     true_states = np.array(true_states)
335
336     # Generate noisy measurements
337     measurements = []
338     for state in true_states:
339         z = h(state)
340         z += np.random.normal(0, np.sqrt(np.diag(ekf.R)))
341         measurements.append(z)
342
343     measurements = np.array(measurements)
344
345     # Run filter
346     estimated_states = []
347     for z in measurements:
348         ekf.predict()
349         ekf.update(z)
350         estimated_states.append(ekf.get_state().x.copy())
351
352     estimated_states = np.array(estimated_states)
353
354     # Plot results
355     plt.figure(figsize=(12, 6))
356
357     plt.subplot(121)
358     plt.plot(true_states[:, 0], true_states[:, 1], 'k-', label='True')
359     plt.plot(estimated_states[:, 0], estimated_states[:, 1], 'r--', label='
Estimated')
360     plt.plot(5.0, 5.0, 'g*', markersize=10, label='Landmark')
361     plt.legend()
362     plt.title('Robot Position Track')
363     plt.xlabel('X')
364     plt.ylabel('Y')
365
366     plt.subplot(122)
367     plt.plot(t, np.unwrap(true_states[:, 2]), 'k-', label='True Heading')
368     plt.plot(t, np.unwrap(estimated_states[:, 2]), 'r--', label='Estimated
Heading')
369     plt.legend()
370     plt.title('Robot Heading')
371     plt.xlabel('Time')
372     plt.ylabel('Heading (rad)')
373
374     plt.tight_layout()
375     plt.show()
376
377 if __name__ == "__main__":
378     print("Running Linear Kalman Filter test...")
379     test_linear_constant_velocity()
380
381     print("\nRunning Extended Kalman Filter test...")
382     test_nonlinear_robot()

```

4.3 Unscented Kalman Filter (UKF)

4.3.1 Sigma points

The Unscented Kalman Filter (UKF) represents a clever solution to the nonlinear filtering problem. At its heart lies the concept of sigma points, which provide an elegant way to handle nonlinear transformations of probability distributions. Let's understand this step by step.

4.3.1.1 The Core Concept

Think of sigma points as carefully chosen “sampling points” that capture the essential statistical properties of your state estimate. Instead of trying to linearize a nonlinear function like the EKF does, the UKF takes these special points and passes them directly through the nonlinear function.

Imagine you're trying to track a robot that moves in a circular path. Rather than attempting to approximate the curved motion with straight lines (like the EKF would), the UKF selects specific points around your current estimate and sees how they each move along the actual curve.

4.3.1.2 How Sigma Points Work

For a state vector of dimension n , we generate $2n + 1$ sigma points. Here's why:

1. One point at the mean of your current state estimate
2. n points “forward” along each principal direction of your uncertainty
3. n points “backward” along those same directions

The mathematical formula for generating these points is:

```
1  $\bar{\mathbf{x}}$  (center point)
2  $\mathbf{X} = \bar{\mathbf{x}} + \sqrt{(\kappa + 1)\mathbf{P}}$  (forward points,  $i = 1, \dots, n$ )
3  $\mathbf{X} = \bar{\mathbf{x}} - \sqrt{(\kappa + 1)\mathbf{P}}$  (backward points,  $i = 1, \dots, n$ )
```

Where:

- $\bar{\mathbf{x}}$ is your current state estimate
- \mathbf{P} is your covariance matrix
- κ is a scaling parameter that helps tune the spread of the points
- $\sqrt{(\kappa + 1)\mathbf{P}}$ represents the matrix square root of $(\kappa + 1)\mathbf{P}$

4.3.1.3 The Weight System

Each sigma point gets assigned two weights:

1. One for recovering the mean (W)
2. One for recovering the covariance (W)

These weights are calculated as:

```
1  $W = 1/(n + \kappa + 1)$  (mean weight for center)
2  $W = 1/(n + \kappa + 1) + (1 - \kappa^2 + \kappa)$  (covariance weight for center)
```

$$W = W = 1/(2(n +)) \quad (\text{weights for other points})$$

4.3.1.4 A Practical Example

Let's consider tracking a robot with a 2D state vector [position, velocity]. You would:

1. Generate 5 sigma points ($2 \times 2 + 1 = 5$)
2. Pass each point through your motion model
3. Use the weighted sum of transformed points to get your new state estimate

The beauty of this approach is that it:

- Captures nonlinear behavior more accurately than EKF
- Doesn't require you to calculate any Jacobian matrices
- Naturally handles strong nonlinearities in your system

4.3.1.5 Parameters and Tuning

Three key parameters control the behavior of sigma points:

- α : Controls spread of points (typically 10^{-3} to 1)
- β : Incorporates prior knowledge ($\beta = 2$ is optimal for Gaussian distributions)
- κ : Secondary scaling parameter (usually 0 or $3-n$)

Think of α as adjusting how far the sigma points spread from the mean - smaller values keep points closer together, while larger values explore more of the state space.

4.3.1.6 Engineering Insights

As an engineering student, here are key points to remember:

1. The number of sigma points scales linearly with state dimension ($2n + 1$), making this method computationally tractable for many real-world problems.
2. The matrix square root computation (usually done via Cholesky decomposition) is the most computationally intensive part.
3. The method preserves the mean and covariance of your distribution up to the third order (Taylor series expansion) for any nonlinear function.

The UKF's sigma point approach often provides better results than the EKF, especially when:

- Your system has strong nonlinearities
- You need better consistency in your uncertainty estimates
- You want to avoid computing Jacobian matrices

4.3.2 Unscented transform

The Unscented Transform is a method for estimating how probability distributions change when they go through nonlinear transformations. Think of it as a sophisticated way to answer the question: "If I have a random variable x with a known distribution, and I put it through a nonlinear function $y = f(x)$, what's the distribution of y ?"

4.3.2.1 Core Principle

The fundamental insight of the UT is that it's easier to approximate a probability distribution than to approximate an arbitrary nonlinear function. Instead of trying to linearize the nonlinear function (as the EKF does), the UT works by:

1. Choosing a set of specific points (sigma points) that capture the key statistical properties of your input distribution
2. Passing these points through your nonlinear function
3. Reconstructing the statistics of the output distribution from the transformed points

4.3.2.2 The Mathematical Framework

Let's say we have a random variable x with mean μ and covariance Σ . The UT proceeds in three key steps:

4.3.2.2.1 Step 1: Sigma Point Generation We generate a set of sigma points $\{X_i\}$ and weights $\{W_i\}$ using the formulas:

$$X_0 = \mu, X_i = \mu + (\sqrt{(n + \lambda)\Sigma})_i, X_{i+n} = \mu - (\sqrt{(n + \lambda)\Sigma})_i$$

Where:

- n is the dimension of x
- λ is a scaling parameter
- $(\sqrt{(n + \lambda)\Sigma})_i$ is the i th column of the matrix square root

4.3.2.2.2 Step 2: Nonlinear Transformation Each sigma point is transformed through the nonlinear function:

$$Y_i = f(X_i)$$

This is where the UT shows its elegance - we're using the actual nonlinear function, not an approximation.

4.3.2.2.3 Step 3: Statistical Recovery The mean and covariance of the transformed distribution are approximated by:

$$\mu_y = \sum_i W_i^m Y_i, \Sigma_y = \sum_i W_i^c (Y_i - \mu_y)(Y_i - \mu_y)^T$$

4.3.2.3 Properties and Insights

The UT has several remarkable properties that make it powerful for engineering applications:

1. **Accuracy:** It captures the mean and covariance accurately to the third order (Taylor series) for any nonlinear function. For Gaussian inputs, it's accurate to the fourth order.
2. **No Derivatives:** Unlike the EKF's linearization, we don't need to calculate any Jacobian matrices. This is particularly valuable when:

- Your function isn't easily differentiable
 - The Jacobian computation is computationally expensive
 - You're working with discontinuous functions
3. Computational Efficiency: For an n -dimensional state, you only need $2n + 1$ sigma points. This makes it practical for real-time applications.

4.3.2.4 A Practical Example

Let's consider a simple nonlinear transformation: converting from polar to Cartesian coordinates.

If we have a point in polar coordinates (r, θ) with some uncertainty, the transformation to Cartesian coordinates is:

$$x = r \cos(\theta) \quad y = r \sin(\theta)$$

The EKF would linearize these equations around the mean. The UT instead:

1. Generates sigma points in polar coordinates
2. Transforms each point using the exact nonlinear equations
3. Reconstructs the Cartesian statistics from the transformed points

This provides a much better approximation of the true uncertainty, especially when the angular uncertainty is large.

4.3.2.5 Implementation Considerations

When implementing the UT, several practical aspects need attention:

1. Parameter Selection:
 - controls the spread of sigma points (typically 10 to 1)
 - incorporates prior knowledge ($\kappa = 2$ for Gaussian distributions)
 - helps optimize higher-order moments (often 0 or $3-n$)
2. Numerical Stability:
 - Use Cholesky decomposition for the matrix square root
 - Monitor the condition number of your covariance matrices
 - Consider using square-root implementations for better numerical properties

4.3.3 Algorithm implementation

Let's create a detailed implementation of an Unscented Kalman Filter in Python, with comprehensive documentation and example usage.

The key components of this implementation:

Core UKF Structure

- **UKFParameters:** Encapsulates the key parameters (κ, α, β) that control the behavior of the unscented transform

- `UnscentedKalmanFilter`: Main filter class implementing the UKF algorithm

Key Methods

- `generate_sigma_points`: Implements the scaled unscented transform to generate sigma points
- `predict`: Propagates sigma points through the process model
- `update`: Incorporates measurements using the unscented transform

Test Scenario:

- Nonlinear tracking problem with range and bearing measurements
- Shows how the UKF handles nonlinear measurement models effectively

```

1 import numpy as np
2 from dataclasses import dataclass
3 from typing import Optional, Tuple, Callable
4 import matplotlib.pyplot as plt
5
6
7 @dataclass
8 class UKFParameters:
9     """
10     Parameters for the Unscented Kalman Filter
11
12     Attributes:
13         alpha: Controls the spread of sigma points (typically 1e-4 to 1)
14         beta: Incorporates prior knowledge (2.0 is optimal for Gaussian)
15         kappa: Secondary scaling parameter (typically 0 or 3-n)
16         n: State dimension
17     """
18     n: int
19     alpha: float = 1e-3
20     beta: float = 2.0
21     kappa: float = 0.0
22
23     @property
24     def lambda_param(self) -> float:
25         """Calculate lambda parameter for sigma point scaling"""
26         return self.alpha**2 * (self.n + self.kappa) - self.n
27
28 class UnscentedKalmanFilter:
29     """
30     Implementation of the Unscented Kalman Filter
31
32     This implementation uses the scaled unscented transform and supports
33     arbitrary nonlinear process and measurement models.
34     """
35
36     def __init__(self, dim_x: int, dim_z: int,
37                  process_fn: Callable, measurement_fn: Callable,
38                  alpha: float = 1e-3, beta: float = 2.0, kappa: float =
39                  0.0):
40         """
41         Initialize the UKF
42
43         Args:

```

```

43         dim_x: State dimension
44         dim_z: Measurement dimension
45         process_fn: State transition function f(x, dt)
46         measurement_fn: Measurement function h(x)
47         alpha: Spread parameter
48         beta: Prior knowledge parameter
49         kappa: Secondary scaling parameter
50     """
51     self.dim_x = dim_x
52     self.dim_z = dim_z
53     self.f = process_fn
54     self.h = measurement_fn
55
56     # Initialize filter parameters
57     self.params = UKFParameters(n=dim_x, alpha=alpha, beta=beta, kappa=
kappa)
58
59     # Initialize state estimate and covariance
60     self.x = np.zeros(dim_x)
61     self.P = np.eye(dim_x)
62
63     # Initialize noise covariances
64     self.Q = np.eye(dim_x) # Process noise
65     self.R = np.eye(dim_z) # Measurement noise
66
67     # Compute number of sigma points and weights once
68     self.n_sigma = 2 * dim_x + 1
69     self._compute_weights()
70
71     def _compute_weights(self) -> None:
72         """Compute weights for sigma points"""
73         n = self.dim_x
74         = self.params.lambda_param
75
76         # Weight for mean at center point
77         self.Wm = np.zeros(self.n_sigma)
78         self.Wm[0] = 1 / (n + )
79
80         # Weight for covariance at center point
81         self.Wc = np.zeros(self.n_sigma)
82         self.Wc[0] = self.Wm[0] + (1 - self.params.alpha**2 + self.params.
beta)
83
84         # Weights for remaining points
85         weight = 1 / (2 * (n + ))
86         self.Wm[1:] = weight
87         self.Wc[1:] = weight
88
89     def generate_sigma_points(self, x: np.ndarray, P: np.ndarray) -> np.
ndarray:
90         """
91         Generate sigma points using the scaled unscented transform
92
93         Args:
94             x: State mean
95             P: State covariance
96
97         Returns:

```



```

98         sigma_points: Array of sigma points (2n+1 x n)
99     """
100     n = self.dim_x
101     = self.params.lambda_param
102
103     # Calculate square root of (n + )P using Cholesky decomposition
104     U = np.linalg.cholesky((n + ) * P)
105
106     # Initialize sigma points matrix
107     sigma_points = np.zeros((self.n_sigma, n))
108     sigma_points[0] = x
109
110     # Generate remaining sigma points
111     for i in range(n):
112         sigma_points[i + 1] = x + U[i]
113         sigma_points[n + i + 1] = x - U[i]
114
115     return sigma_points
116
117 def predict(self, dt: float) -> None:
118     """
119     Predict step of the UKF
120
121     Args:
122         dt: Time step
123     """
124     # Generate sigma points
125     sigma_points = self.generate_sigma_points(self.x, self.P)
126
127     # Transform sigma points through process model
128     transformed_sigmas = np.array([self.f(sigma, dt) for sigma in
sigma_points])
129
130     # Recover mean and covariance
131     self.x = np.sum(self.Wm.reshape(-1, 1) * transformed_sigmas, axis
=0)
132
133     # Calculate covariance of prediction
134     self.P = np.zeros_like(self.P)
135     for i in range(self.n_sigma):
136         diff = (transformed_sigmas[i] - self.x).reshape(-1, 1)
137         self.P += self.Wc[i] * diff @ diff.T
138
139     # Add process noise
140     self.P += self.Q
141
142 def update(self, z: np.ndarray) -> None:
143     """
144     Update step of the UKF
145
146     Args:
147         z: Measurement vector
148     """
149     # Generate sigma points
150     sigma_points = self.generate_sigma_points(self.x, self.P)
151
152     # Transform sigma points through measurement model
153     transformed_sigmas = np.array([self.h(sigma) for sigma in

```

```

sigma_points])

154
155     # Predict measurement
156     z_pred = np.sum(self.Wm.reshape(-1, 1) * transformed_sigmas, axis
=0)

157
158     # Calculate innovation covariance
159     S = np.zeros((self.dim_z, self.dim_z))
160     Pxz = np.zeros((self.dim_x, self.dim_z))
161
162     for i in range(self.n_sigma):
163         diff_z = (transformed_sigmas[i] - z_pred).reshape(-1, 1)
164         diff_x = (sigma_points[i] - self.x).reshape(-1, 1)
165
166         S += self.Wc[i] * diff_z @ diff_z.T
167         Pxz += self.Wc[i] * diff_x @ diff_z.T
168
169     # Add measurement noise
170     S += self.R
171
172     # Calculate Kalman gain
173     K = Pxz @ np.linalg.inv(S)
174
175     # Update state and covariance
176     self.x = self.x + K @ (z - z_pred)
177     self.P = self.P - K @ S @ K.T
178
179 def test_nonlinear_tracking():
180     """
181     Test the UKF with a nonlinear tracking problem
182     """
183     def process_fn(x, dt):
184         """Nonlinear process model: constant turn rate and velocity"""
185         F = np.array([[1, 0, dt, 0],
186                       [0, 1, 0, dt],
187                       [0, 0, 1, 0],
188                       [0, 0, 0, 1]])
189         return F @ x
190
191     def measurement_fn(x):
192         """Nonlinear measurement model: range and bearing"""
193         px, py = x[0], x[1]
194         r = np.sqrt(px**2 + py**2)
195         theta = np.arctan2(py, px)
196         return np.array([r, theta])
197
198     # Initialize UKF
199     ukf = UnscentedKalmanFilter(
200         dim_x=4, # [px, py, vx, vy]
201         dim_z=2, # [range, bearing]
202         process_fn=process_fn,
203         measurement_fn=measurement_fn
204     )
205
206     # Set initial state and covariance
207     ukf.x = np.array([0., 0., 1., 1.])
208     ukf.P *= 0.1
209

```

```

210 # Set noise covariances
211 ukf.Q = np.eye(4) * 0.1 # Process noise
212 ukf.R = np.array([[0.1, 0], # Measurement noise
213                  [0, 0.01]])
214
215 # Generate true trajectory and measurements
216 dt = 0.1
217 t = np.arange(0, 10, dt)
218 n_steps = len(t)
219
220 true_states = np.zeros((n_steps, 4))
221 measurements = np.zeros((n_steps, 2))
222 estimated_states = np.zeros((n_steps, 4))
223
224 x_true = ukf.x.copy()
225
226 for k in range(n_steps):
227     # Generate true trajectory
228     x_true = process_fn(x_true, dt) + np.random.multivariate_normal(
229         np.zeros(4), ukf.Q*0.1)
230     true_states[k] = x_true
231
232     # Generate noisy measurement
233     z_true = measurement_fn(x_true)
234     measurements[k] = z_true + np.random.multivariate_normal(
235         np.zeros(2), ukf.R)
236
237     # UKF prediction and update
238     ukf.predict(dt)
239     ukf.update(measurements[k])
240     estimated_states[k] = ukf.x
241
242 # Plot results
243 plt.figure(figsize=(15, 5))
244
245 # Position track
246 plt.subplot(121)
247 plt.plot(true_states[:, 0], true_states[:, 1], 'b-', label='True')
248 plt.plot(estimated_states[:, 0], estimated_states[:, 1], 'r--',
249         label='Estimated')
250 plt.scatter(measurements[:, 0] * np.cos(measurements[:, 1]),
251            measurements[:, 0] * np.sin(measurements[:, 1]),
252            c='g', alpha=0.2, label='Measurements')
253 plt.legend()
254 plt.title('Position Track')
255 plt.xlabel('X Position')
256 plt.ylabel('Y Position')
257 plt.grid(True)
258
259 # Velocity estimates
260 plt.subplot(122)
261 plt.plot(t, true_states[:, 2], 'b-', label='True Vx')
262 plt.plot(t, true_states[:, 3], 'b--', label='True Vy')
263 plt.plot(t, estimated_states[:, 2], 'r-', label='Est Vx')
264 plt.plot(t, estimated_states[:, 3], 'r--', label='Est Vy')
265 plt.legend()
266 plt.title('Velocity Estimates')
267 plt.xlabel('Time')

```

```
268     plt.ylabel('Velocity')
269     plt.grid(True)
270
271     plt.tight_layout()
272     plt.show()
273
274 if __name__ == "__main__":
275     test_nonlinear_tracking()
```

Chapter 5

Particle Filtering

- Particle Filtering
 - Monte Carlo methods
 - Importance sampling
 - Particle representation
 - Sequential Importance Sampling (SIS)
 - Resampling techniques
 - Sample degeneracy and impoverishment
 - Adaptive particle filtering

Chapter 6

Advanced Topics

- ☐ Advanced Topics
 - ☐ Multi-hypothesis tracking
 - ☐ SLAM basics
 - ☐ Sensor fusion techniques
 - ☐ Loop closure
 - ☐ Global vs local localization

Chapter 7

MATLAB Implementation: Kalman Filter

- ☐ MATLAB Implementation: Kalman Filter
 - ☐ Linear KF implementation
 - ☐ State prediction
 - ☐ Measurement update
 - ☐ Covariance propagation
 - ☐ EKF implementation
 - ☐ System modeling
 - ☐ Jacobian computation
 - ☐ Filter implementation
 - ☐ Visualization and analysis
 - ☐ Performance evaluation

Chapter 8

MATLAB Implementation: Particle Filter

- ☐ MATLAB Implementation: Particle Filter
 - ☐ Basic PF framework
 - ☐ Particle initialization
 - ☐ Motion model implementation
 - ☐ Measurement model
 - ☐ Weight computation
 - ☐ Resampling implementation
 - ☐ Visualization tools
 - ☐ Performance metrics

Chapter 9

Practical Applications

- Practical Applications
 - Vehicle localization case studies
 - Robot navigation examples
 - Integration with mapping
 - Real-world challenges
 - Best practices and optimization

Chapter 10

Project Work

- ☐ Project Work
 - ☐ Implementation exercises
 - ☐ Real dataset analysis
 - ☐ Performance comparison of different filters
 - ☐ Parameter tuning
 - ☐ Documentation and presentation